



# Real-time Data Pipelines with

# Structured Streaming in Spark™

Tathagata “TD” Das

 @tathadas

DataEngConf 2018

18<sup>th</sup> April, San Francisco



# About Me

Started Spark Streaming project in AMPLab, UC Berkeley

Currently focused on building Structured Streaming

PMC Member of  Apache Spark™

Engineer on the StreamTeam @  databricks®  
*"we make all your streams come true"*

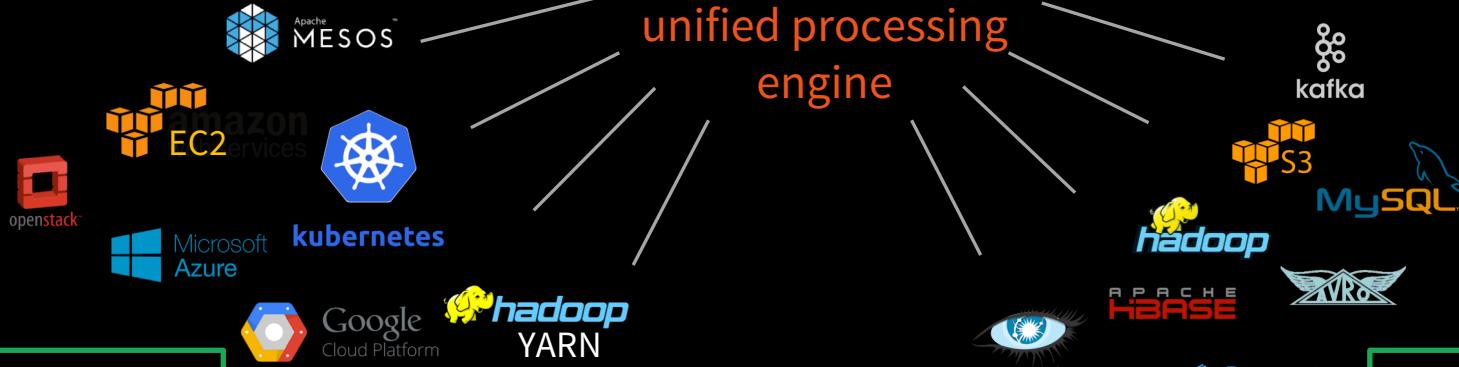
# Applications



Streaming   SQL   ML   Graph



unified processing  
engine

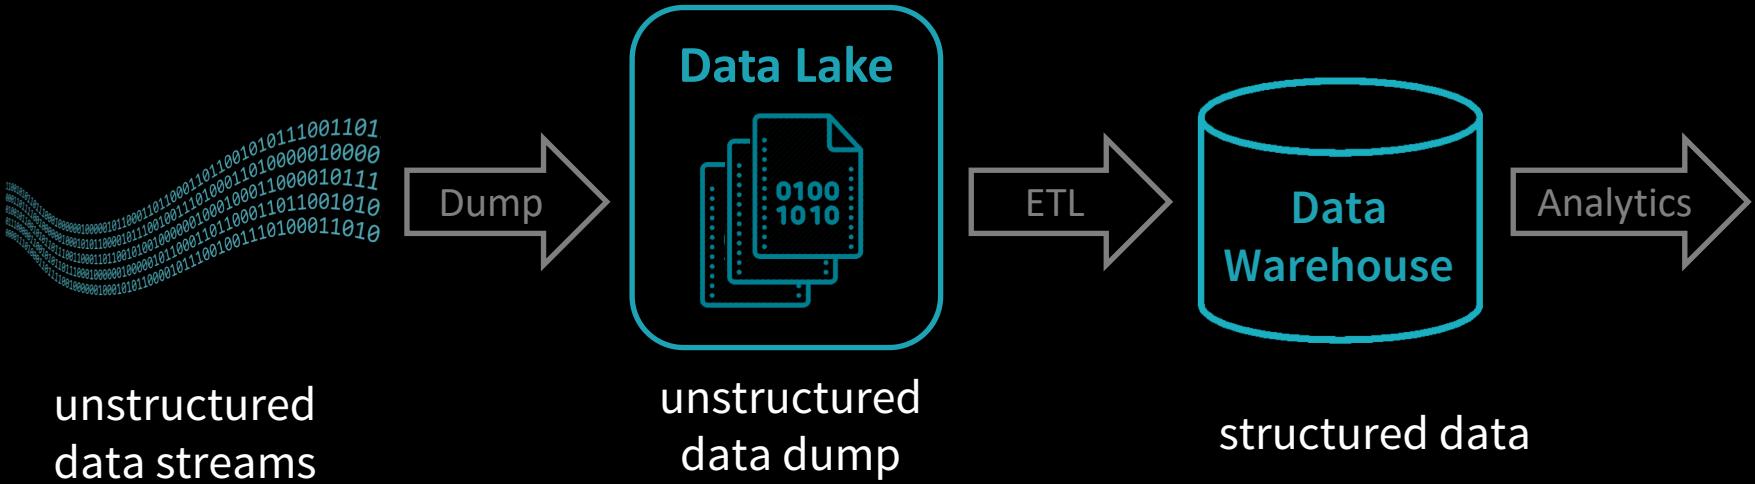


Environments

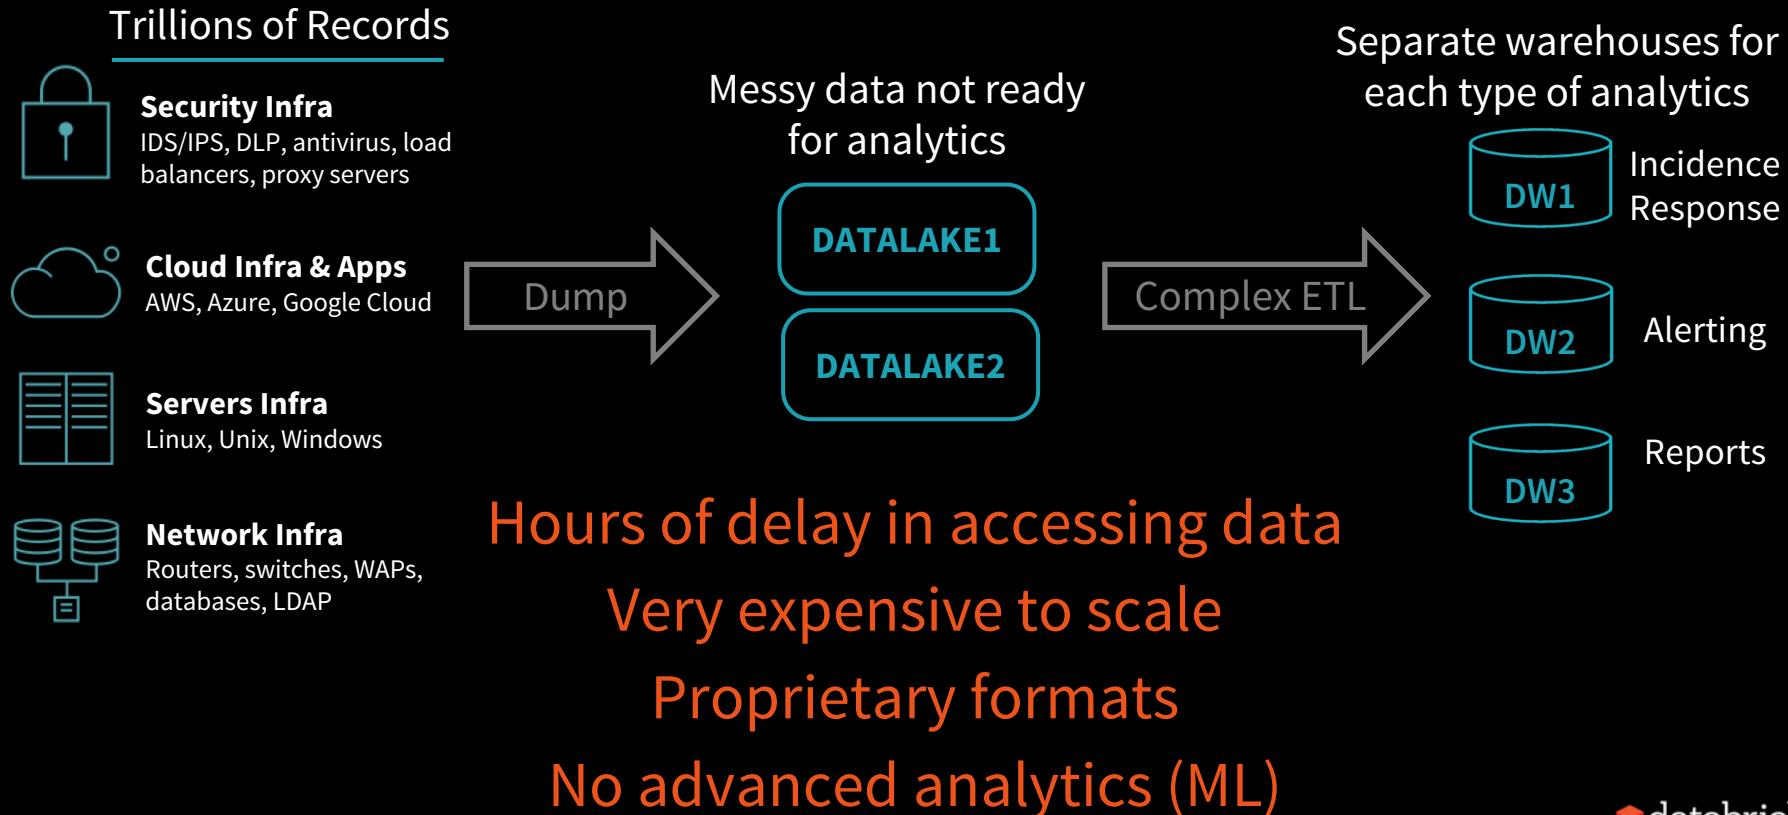
Data Sources

databricks

# Data Pipelines – 10000ft view



# Data Pipeline @ Fortune 100 Company



# New Pipeline @ Fortune 100 Company



Data usable in minutes/seconds

Easy to scale

Open formats

Enables advanced analytics



# STRUCTURED STREAMING

you  
should not have to  
reason about streaming

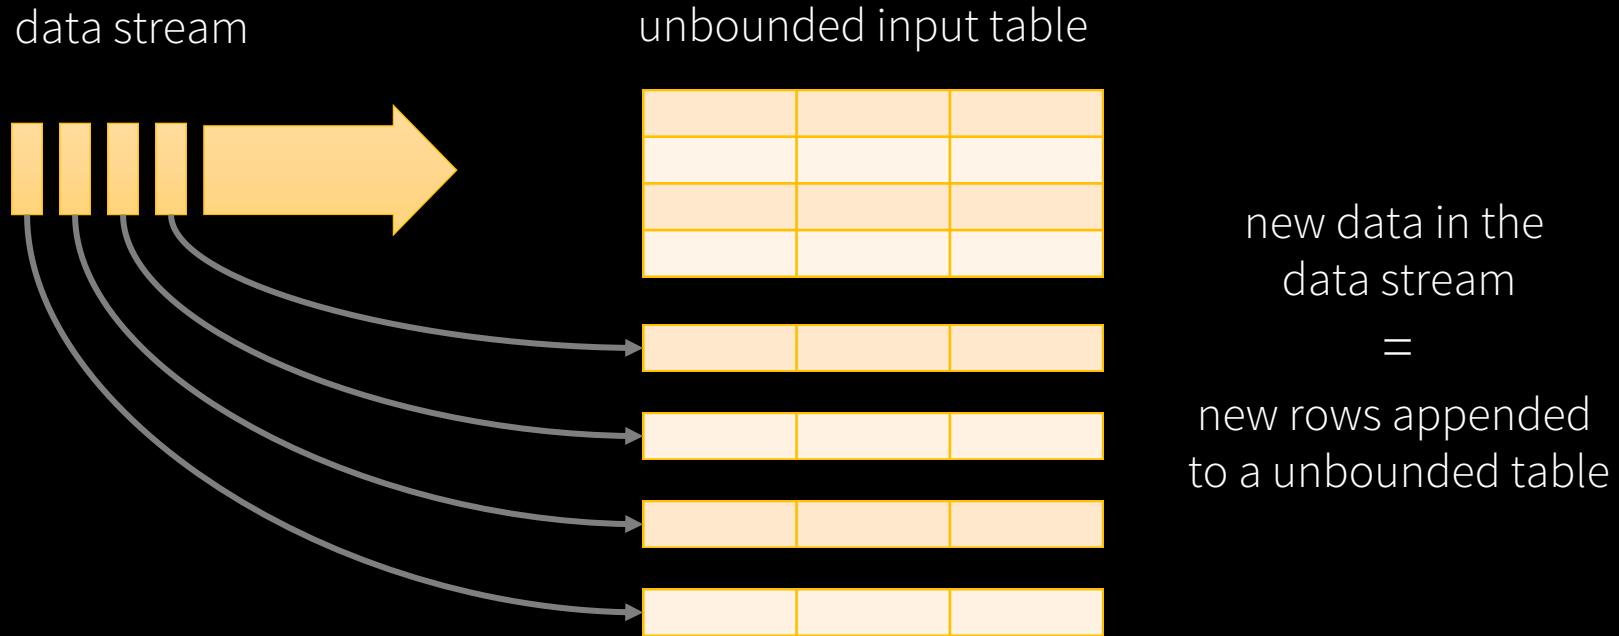
you  
should write simple queries

&

Spark

should continuously update the answer

# Treat Streams as Unbounded Tables



# Anatomy of a Streaming Query

## Example

Read JSON data from Kafka

Parse nested JSON

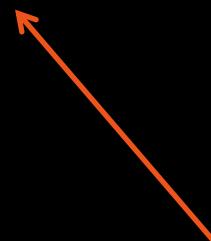
Store in structured Parquet table

Get end-to-end failure guarantees



# Anatomy of a Streaming Query

```
spark.readStream.format("kafka")  
  .option("kafka.bootstrap.servers",...)  
  .option("subscribe", "topic")  
  .load()
```



returns a  
DataFrame



## Source

Specify where to read data from

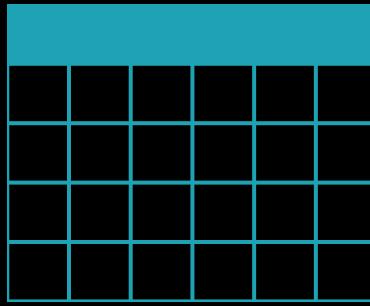
Built-in support for Files / Kafka / Kinesis\*

Can include multiple sources of different types using `join()` / `union()`

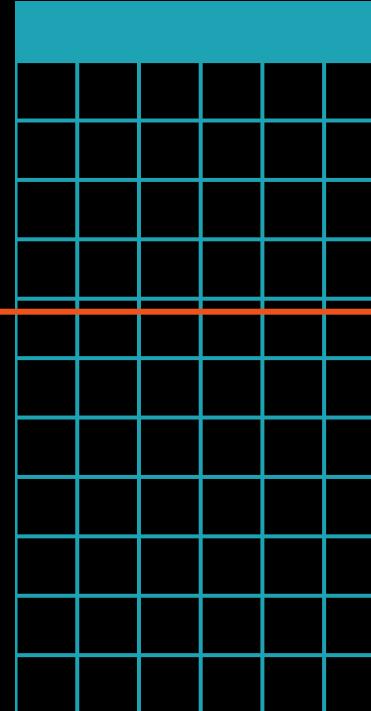
\*Available only on [Databricks Runtime](#)

# DataFrame $\leftrightarrow$ Table

static data =  
bounded table



streaming data =  
unbounded table



Single  
API !

# DataFrame/Dataset

## SQL

```
spark.sql(  
    "SELECT type, sum(signal)  
    FROM devices  
    GROUP BY type  
)
```

Most familiar to BI Analysts  
Supports SQL-2003, HiveQL

## DataFrame



```
val df: DataFrame =  
    spark.table("device-data")  
        .groupBy("type")  
        .sum("signal")
```

Great for Data Scientists familiar  
with Pandas, R Dataframes

## Dataset



```
val ds: Dataset[(String, Double)] =  
    spark.table("device-data")  
        .as[DeviceData]  
        .groupByKey(_.type)  
        .mapValues(_.signal)  
        .reduceGroups(_ + _)
```

Great for Data Engineers who  
want compile-time type safety

Same semantics, same performance

Choose your hammer for whatever nail you have!

# Anatomy of a Streaming Query

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
```



Kafka DataFrame

key	value	topic	partition	offset	timestamp
[binary]	[binary]	"topic"	0	345	1486087873
[binary]	[binary]	"topic"	3	2890	1486086721

# Anatomy of a Streaming Query

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
```

## Transformations

Cast bytes from Kafka records to a string, parse it as a json, and generate nested columns

100s of built-in, optimized SQL functions like `from_json`

user-defined functions, lambdas, function literals with `map`, `flatMap`...

# Anatomy of a Streaming Query

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .writeStream
  .format("parquet")
  .option("path", "/parquetTable/") }
```

## Sink

Write transformed output to external storage systems

Built-in support for Files / Kafka

Use **foreach** to execute arbitrary code with the output data

Some sinks are transactional and exactly once (e.g. files)

# Anatomy of a Streaming Query

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .writeStream
  .format("parquet")
  .option("path", "/parquetTable/")
  .trigger("1 minute")
  .option("checkpointLocation", "...")
  .start()
```

## Processing Details

**Trigger:** when to process data

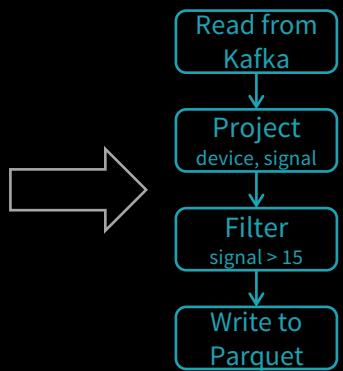
- Fixed interval micro-batches
- As fast as possible micro-batches
- Continuously (new in Spark 2.3)

**Checkpoint location:** for tracking the progress of the query

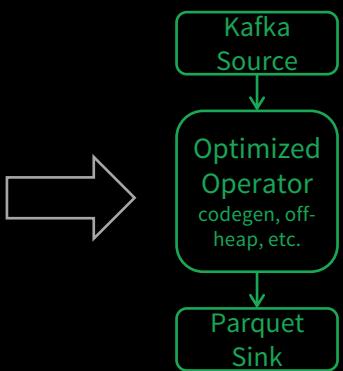
# Spark automatically streamifies!

```
spark.readStream.format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .writeStream
  .format("parquet")
  .option("path", "/parquetTable/")
  .trigger("1 minute")
  .option("checkpointLocation", "...")  
.start()
```

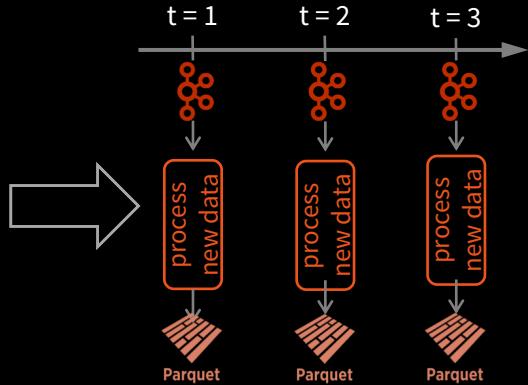
DataFrames,  
Datasets, SQL



Logical  
Plan



Optimized  
Plan



Series of Incremental  
Execution Plans

Spark SQL converts batch-like query to a series of incremental execution plans operating on new batches of data

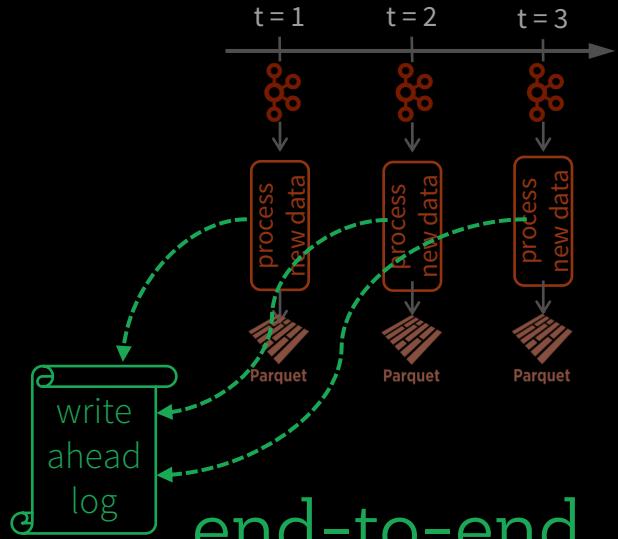
# Fault-tolerance with Checkpointing

## Checkpointing

Saves processed offset info to stable storage  
Saved as JSON for forward-compatibility

Allows recovery from any failure

Can resume after limited changes to your streaming transformations (e.g. adding new filters to drop corrupted data, etc.)



end-to-end  
exactly-once  
guarantees

# Anatomy of a Streaming Query

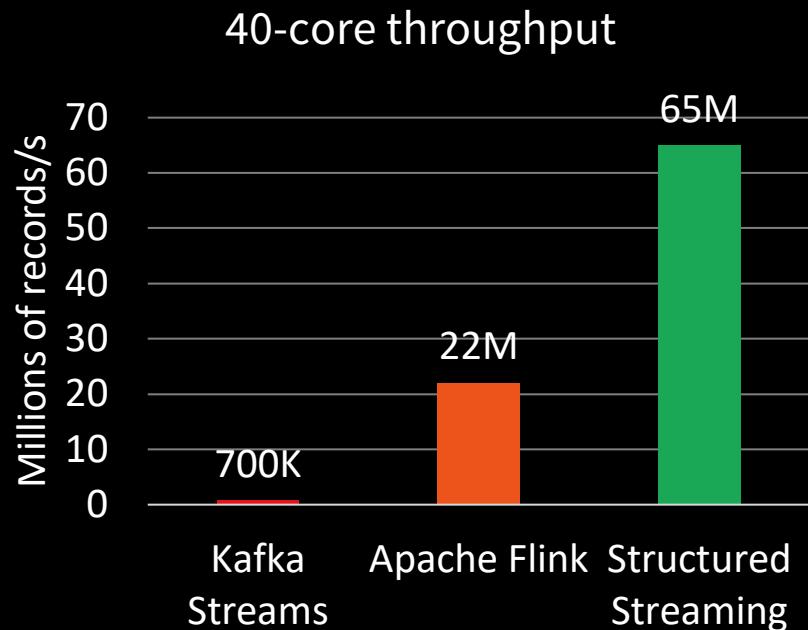
```
spark.readStream.format("kafka")
.option("kafka.bootstrap.servers",...)
.option("subscribe", "topic")
.load()
.selectExpr("cast (value as string) as json")
.select(from_json("json", schema).as("data"))
.writeStream
.format("parquet")
.option("path", "/parquetTable/")
.trigger("1 minute")
.option("checkpointLocation", ...)
.start()
```



Raw data from Kafka available  
as structured data in seconds,  
ready for querying

# Performance: YAHOO! Benchmark

Structured Streaming reuses  
the **Spark SQL Optimizer**  
and **Tungsten Engine**



More details in our [blog post](#)

# Business Logic independent of Execution Mode

```
spark.readStream.format("kafka")
    .option("kafka.bootstrap.servers",...)
    .option("subscribe", "topic")
    .load()
    .selectExpr("cast (value as string) as json")
    .select(from_json("json", schema).as("data"))
    .writeStream
    .format("parquet")
    .option("path", "/parquetTable/")
    .trigger("1 minute")
    .option("checkpointLocation", "...")
    .start()
```



Business logic

# Business Logic independent of Execution Mode

```
spark.read.format("kafka")  
  .option("kafka.bootstrap.servers", ...)  
  .option("subscribe", "topic")  
  .load()
```

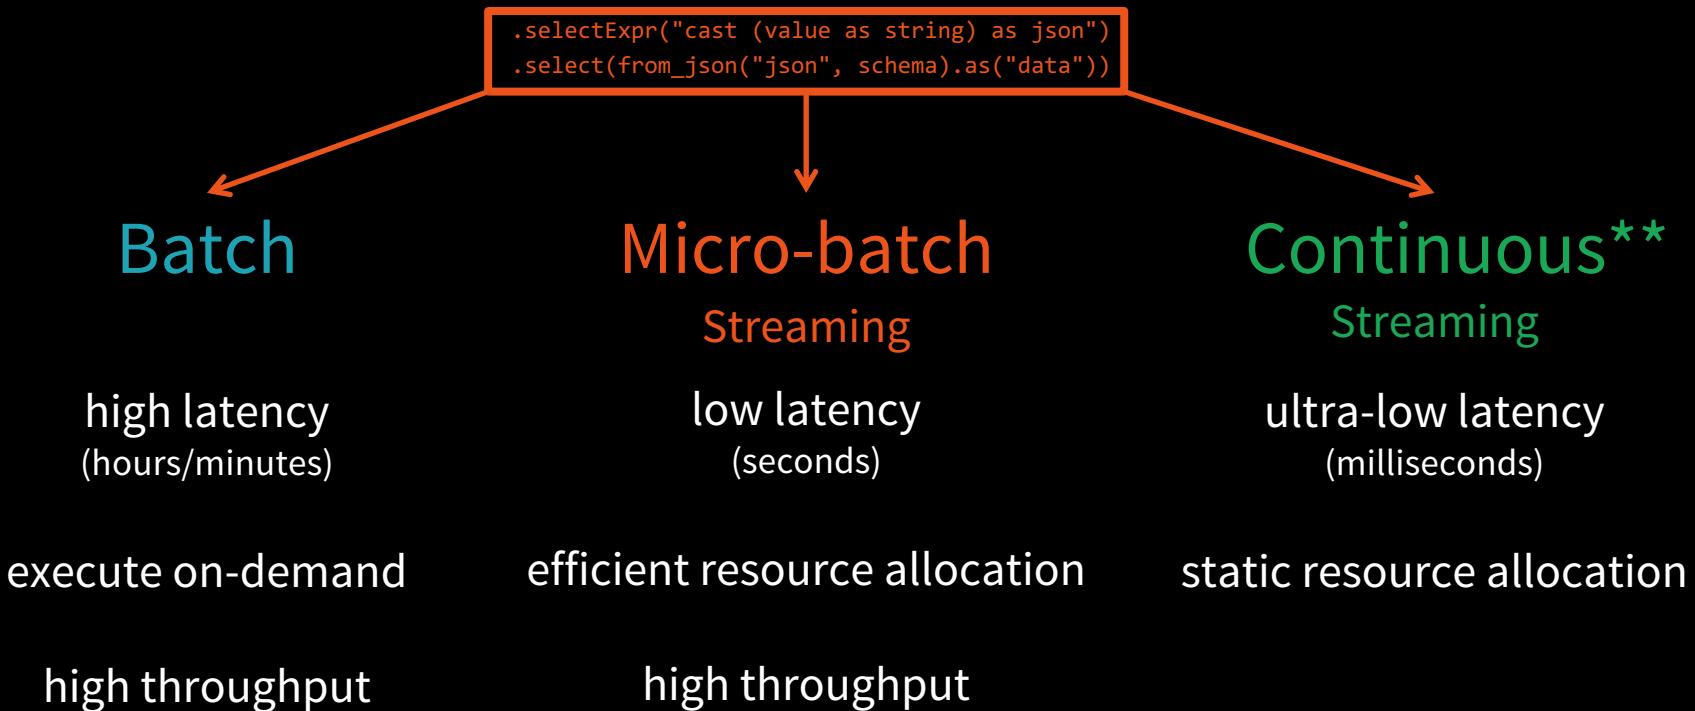
```
  .selectExpr("cast (value as string) as json")  
  .select(from_json("json", schema).as("data"))
```

```
  .write  
  .format("parquet")  
  .option("path", "/parquetTable/")  
  .load()
```

Business logic  
remains unchanged

Peripheral code decides whether  
it's a batch or a streaming query

# Business Logic independent of Execution Mode



\*\*experimental release in Spark 2.3, read our [blog](#)

# Event time Aggregations

Windowing is just another type of grouping in Struct. Streaming

number of records every hour

```
parsedData  
    .groupBy(window("timestamp", "1 hour"))  
    .count()
```

avg signal strength of each device every 10 mins

```
parsedData  
    .groupBy(  
        "device",  
        window("timestamp", "10 mins"))  
    .avg("signal")
```

Support UDAFs!

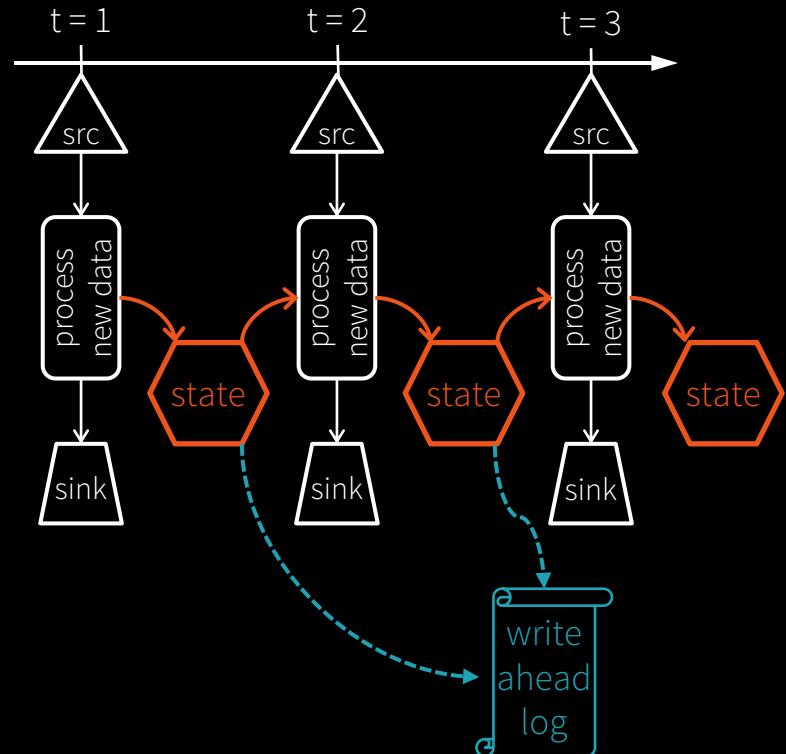
# Stateful Processing for Aggregations

Aggregates has to be saved as  
**distributed state** between triggers

Each trigger reads previous state and  
writes updated state

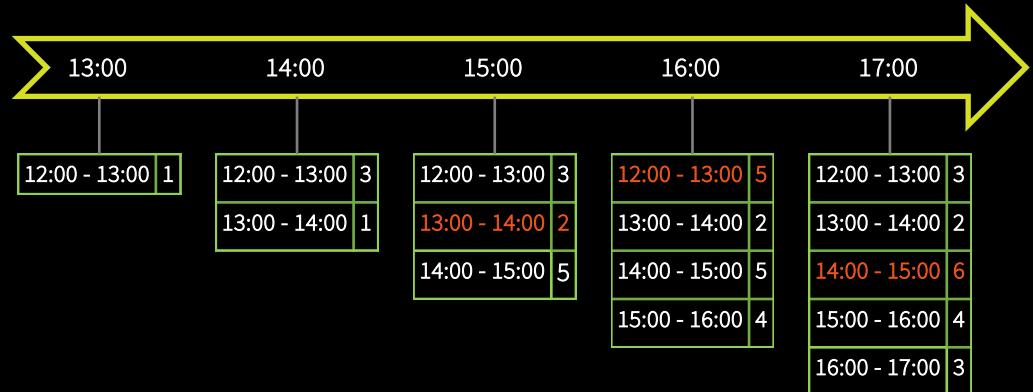
State stored in memory,  
backed by *write ahead log* in HDFS

Fault-tolerant, **exactly-once guarantee!**



# Automatically handles Late Data

Keeping state allows late data to update counts of old windows



But size of the state increases indefinitely if old windows are not dropped

red = state updated with late data

# Watermarking

Watermark - moving threshold of how late data is expected to be and when to drop old state

Trails behind max event time seen by the engine

Watermark delay = trailing gap



# Watermarking

Data newer than watermark may be late, but allowed to aggregate

Data older than watermark is "too late" and dropped

Windows older than watermark automatically deleted to limit the amount of intermediate state

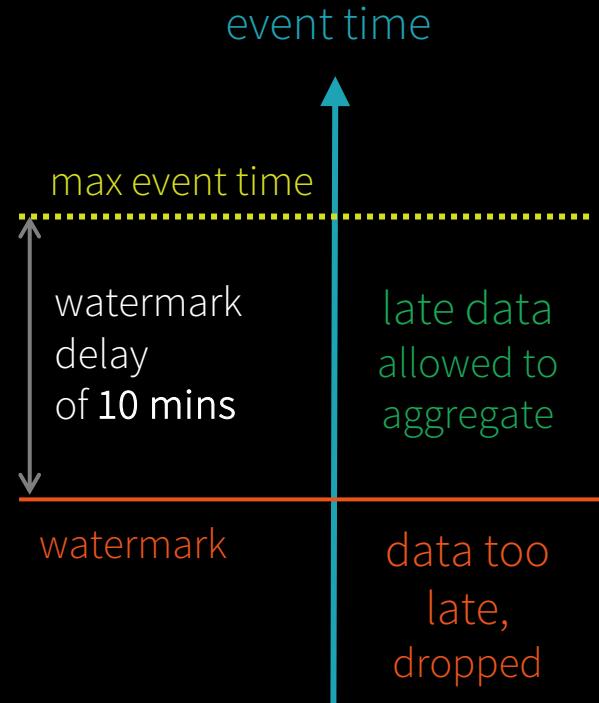


# Watermarking

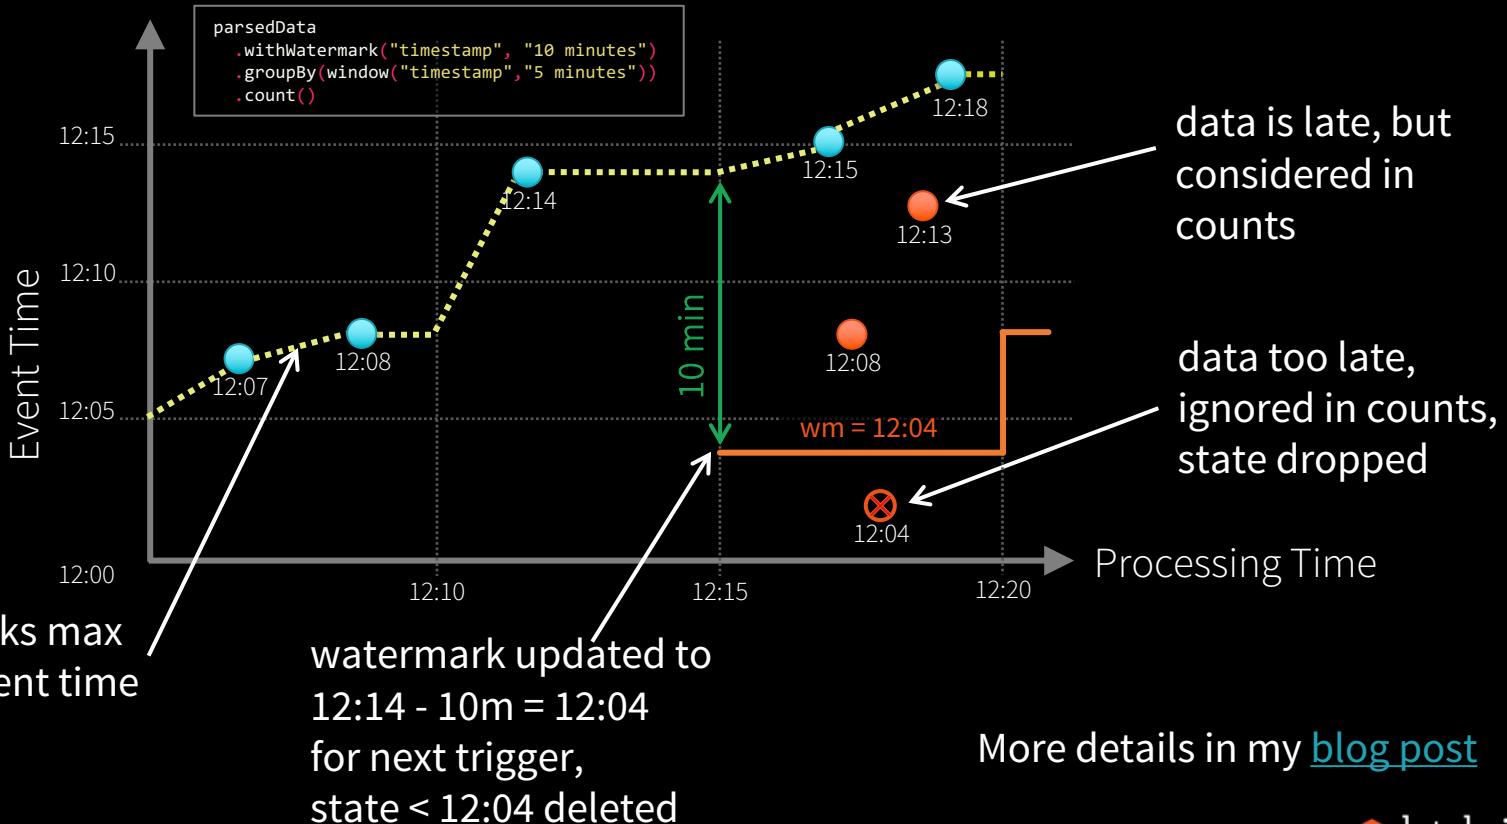
```
parsedData  
  .withWatermark("timestamp", "10 minutes")  
  .groupBy(window("timestamp","5 minutes"))  
  .count()
```

Useful only in stateful operations

Ignored in non-stateful streaming queries and batch queries



# Watermarking



# Other Interesting Operations

## Streaming Deduplication

```
parsedData.dropDuplicates("eventId")
```

## Joins

Stream-batch joins

```
stream1.join(stream2, "device")
```

Stream-stream joins

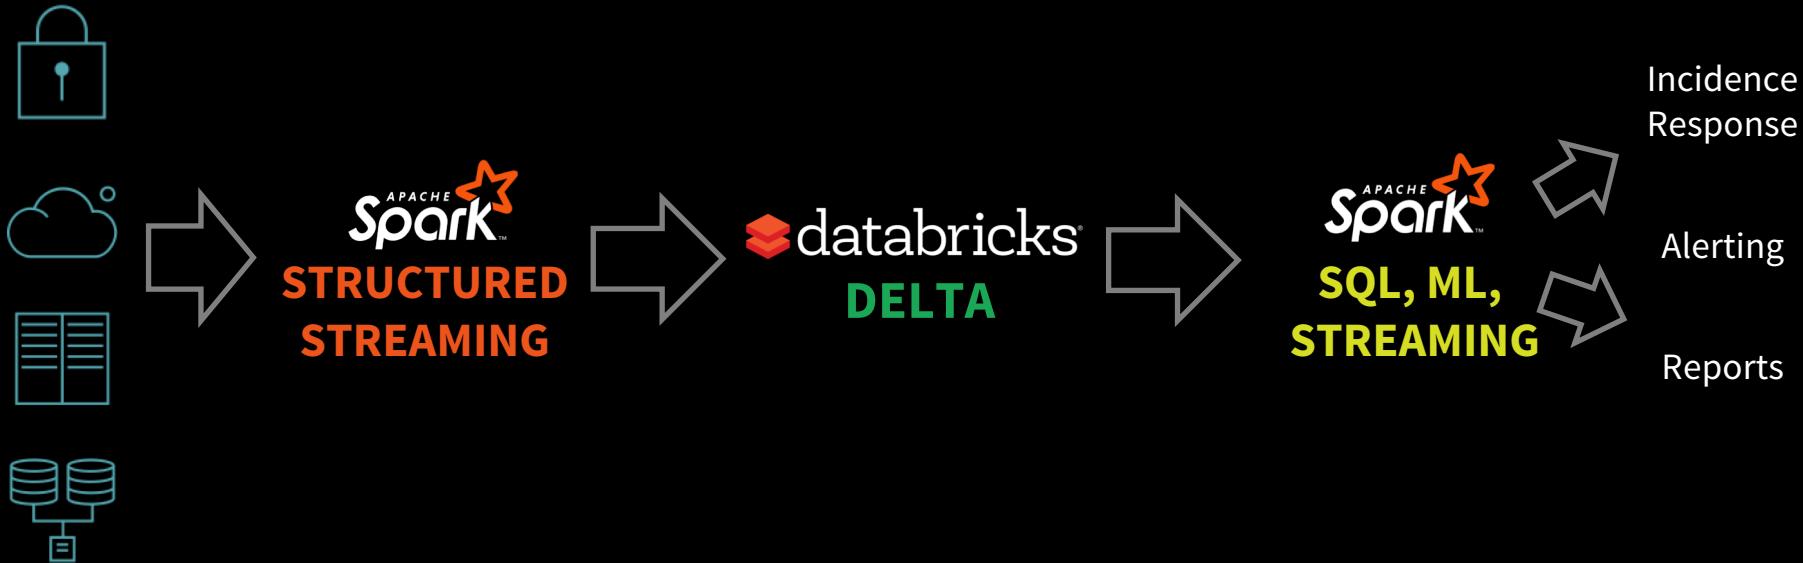
## Arbitrary Stateful Processing

[map|flatMap]GroupsWithState

```
ds.groupByKey(_.id)
  .mapGroupsWithState
    (timeoutConf)
    (mappingWithStateFunc)
```

See my previous Spark Summit [talk](#) and blog posts ([here](#) and [here](#))

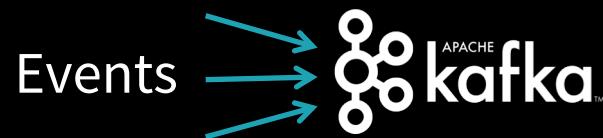
# Data Pipeline with



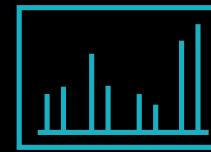


ETL @  databricks®

# Evolution of a Cutting-Edge Data Pipeline



Data Lake

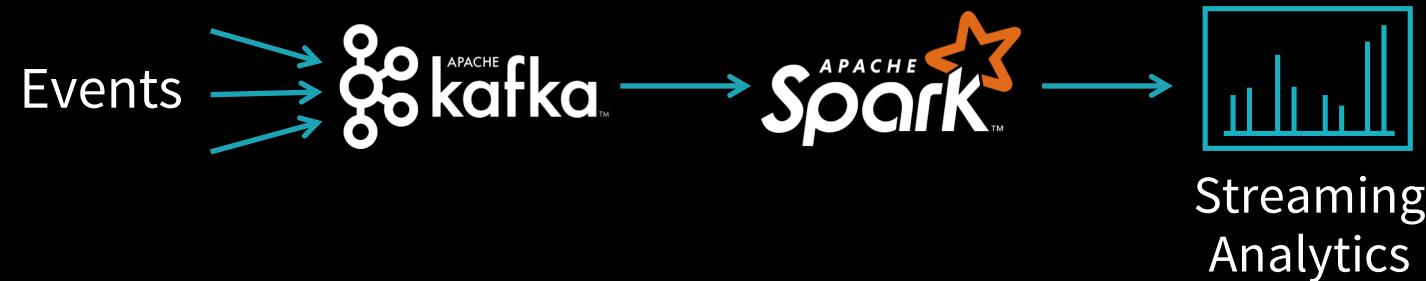


Streaming  
Analytics



Reporting

# Evolution of a Cutting-Edge Data Pipeline

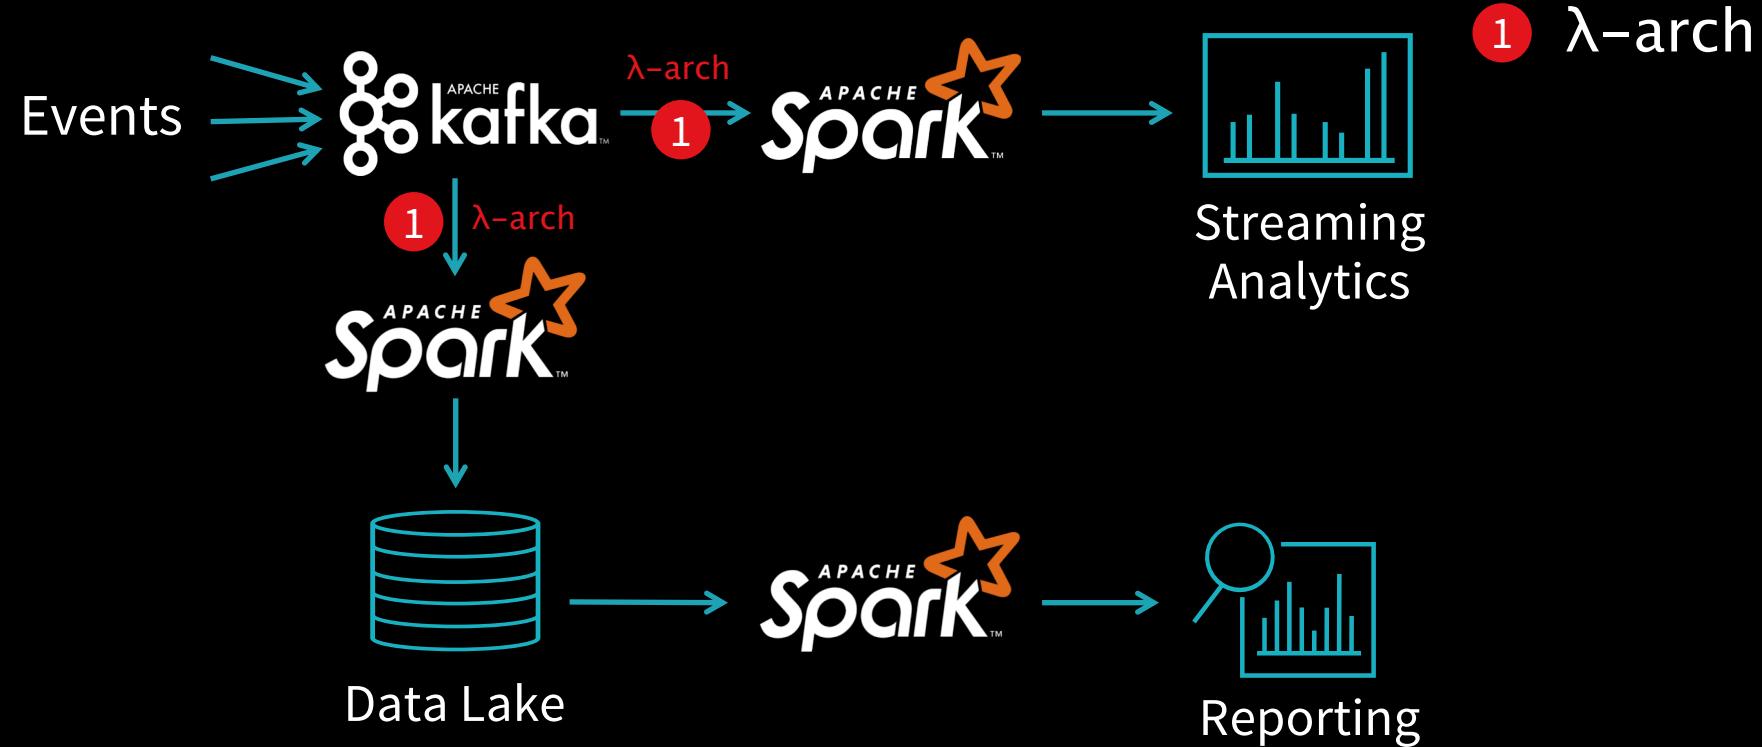


Data Lake

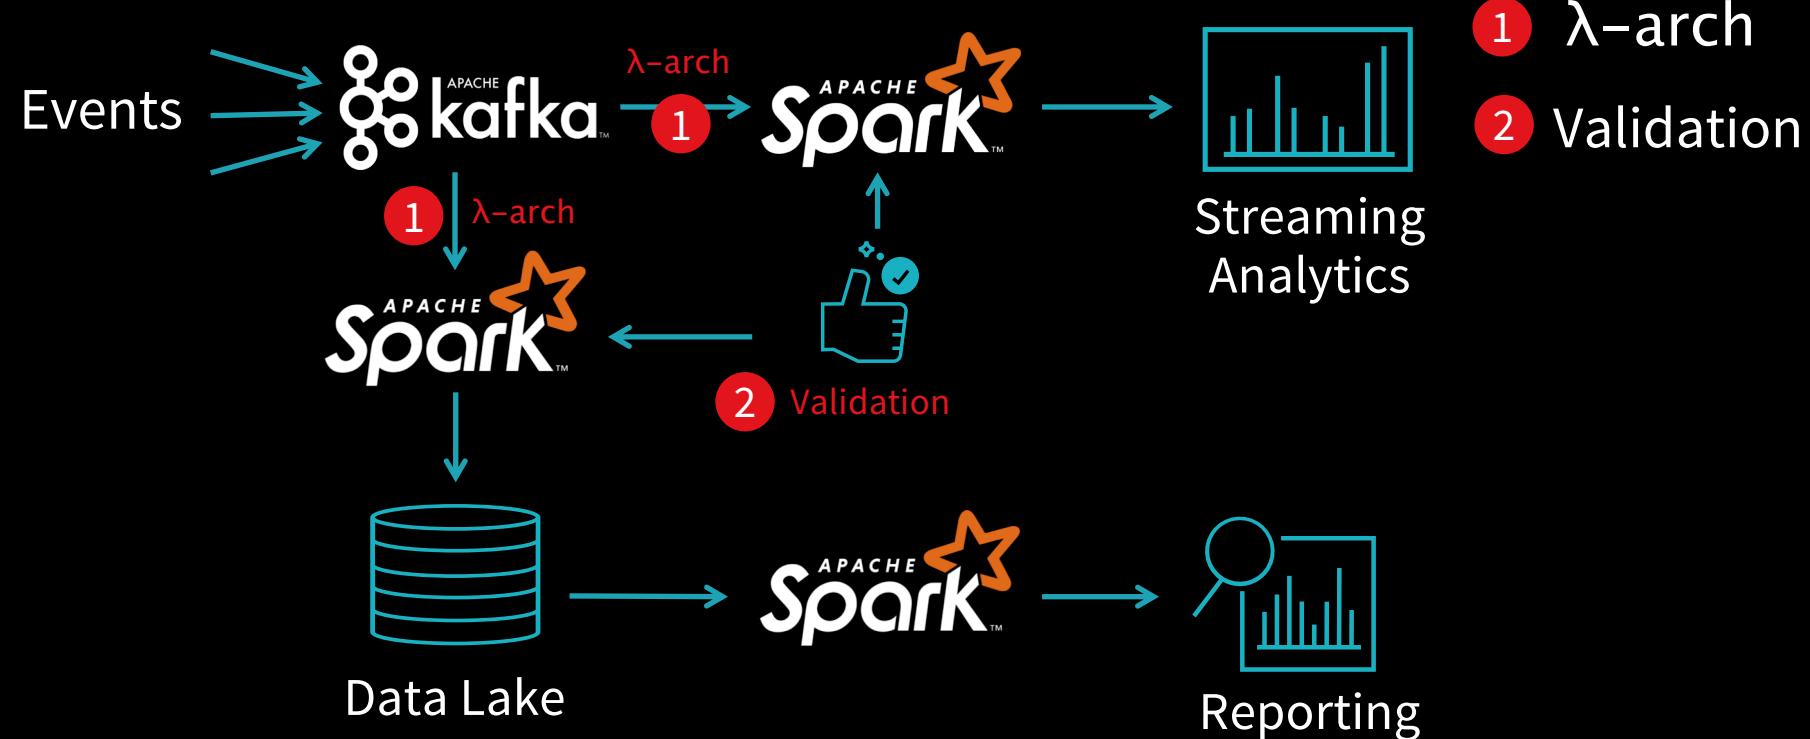


Reporting

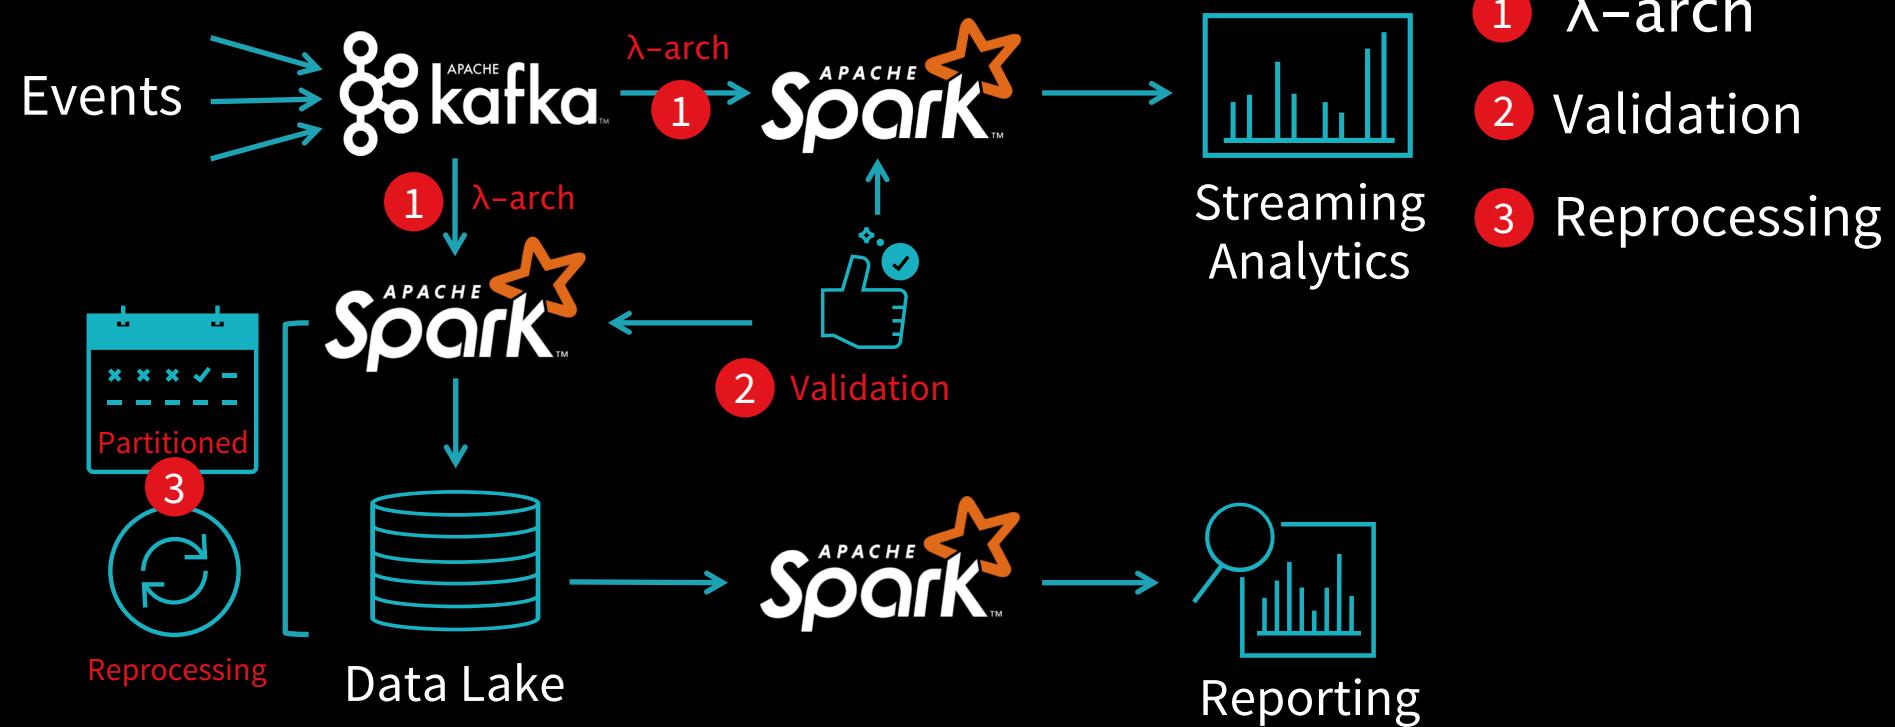
# Challenge #1: Historical Queries?



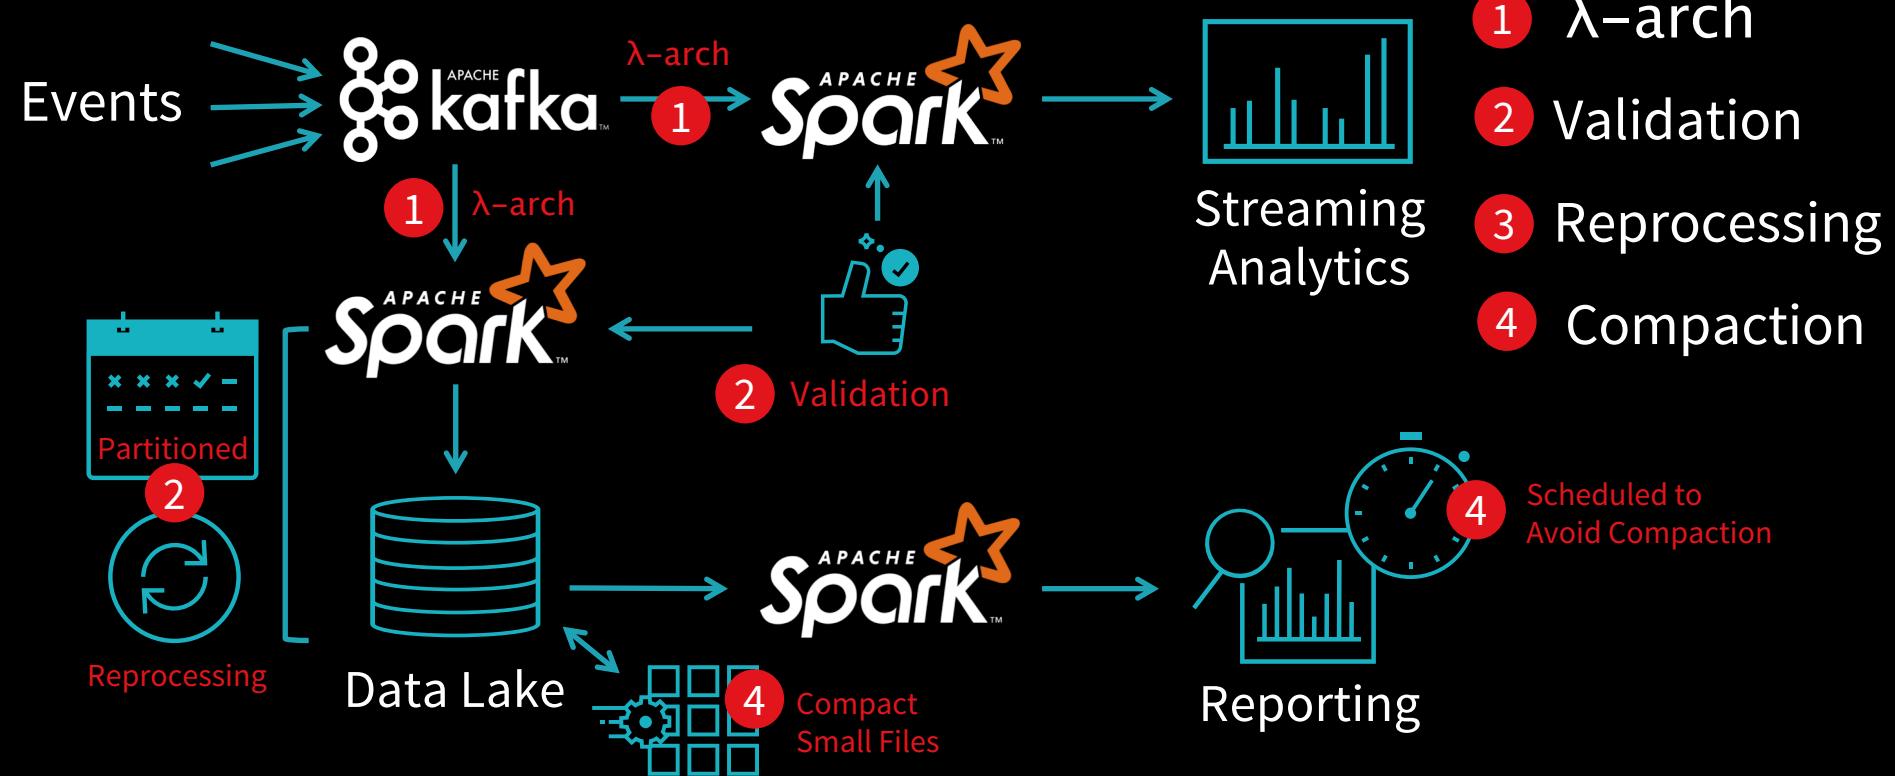
# Challenge #2: Messy Data?



# Challenge #3: Mistakes and Failures?



# Challenge #4: Query Performance?



Let's try it instead with





# databricks<sup>®</sup>

## DELTA

The  
**SCALE**  
of data lake

The  
**RELIABILITY &  
PERFORMANCE**  
of data warehouse

The  
**LOW-LATENCY**  
of streaming

## **THE GOOD OF DATA WAREHOUSES**

- Pristine Data
- Transactional Reliability
- Fast Queries

## **THE GOOD OF DATA LAKES**

- Massive scale on cloud storage
- Open Formats (Parquet, ORC)
- Predictions (ML) & Streaming

# Databricks Delta Combines the Best

**MASSIVE SCALE**

Decouple Compute & Storage

**RELIABILITY**

ACID Transactions & Data Validation

**PERFORMANCE**

Data Indexing & Caching (10-100x)

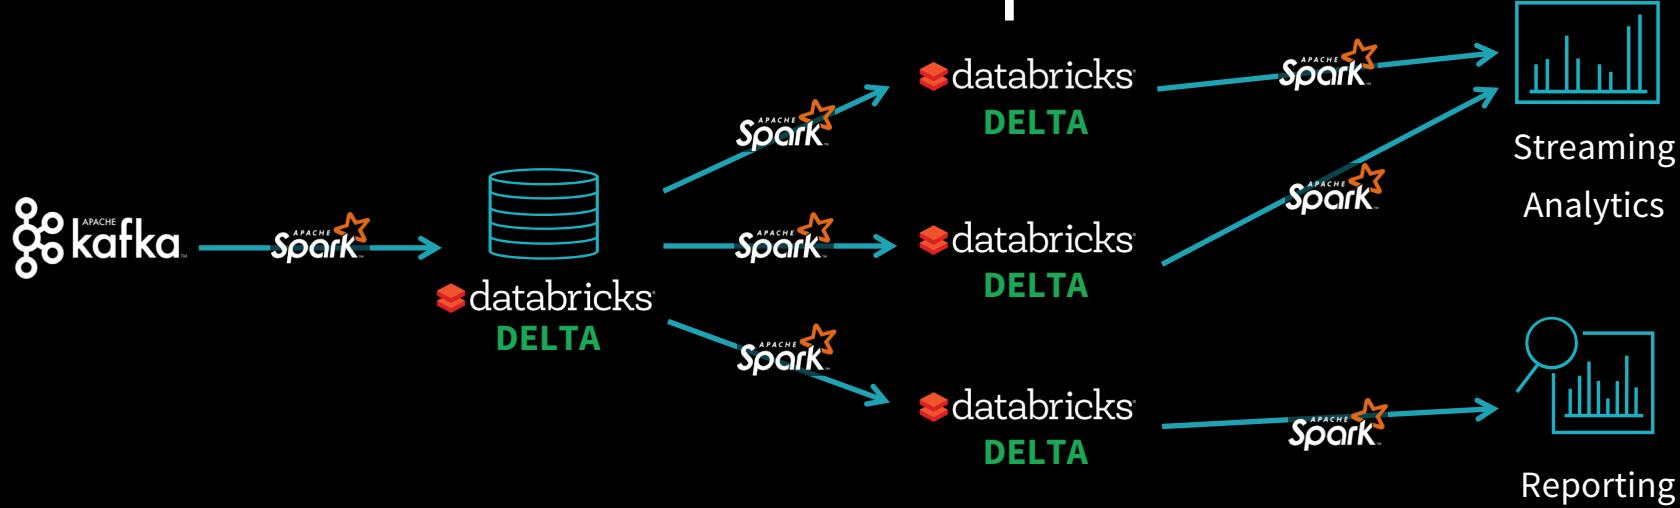
**OPEN**

Data stored as Parquet, ORC, etc.

**LOW-LATENCY**

Integrated with Structured Streaming

# The Canonical Data Pipeline



- ① ~~λ arch~~ → Not needed, Delta handles both short and long term data
- ② Validation ✓ → Easy as data in short term and long term data in one location
- ③ Reprocessing ✓ → Easy and seamless with Delta's transactional guarantees
- ④ Compaction ✓ →

# Accelerate Innovation with Databricks

Increases Data Science Productivity by 5x



## Unified Analytics Platform

### DATABRICKS COLLABORATIVE NOTEBOOKS

Explore Data ➔ Train Models ➔ Serve Models



Open Extensible API's



Improves Performance by 10-20X over Apache Spark



### DATABRICKS RUNTIME

I/O Performance

### DATABRICKS DELTA

Performance Reliability



Higher Performance & Reliability for your Data Lake



Removes Devops & Infrastructure Complexity



### DATABRICKS MANAGED SERVICE

Serverless SLA's



Azure

Databricks Enterprise Security



IoT / STREAMING DATA



CLOUD STORAGE



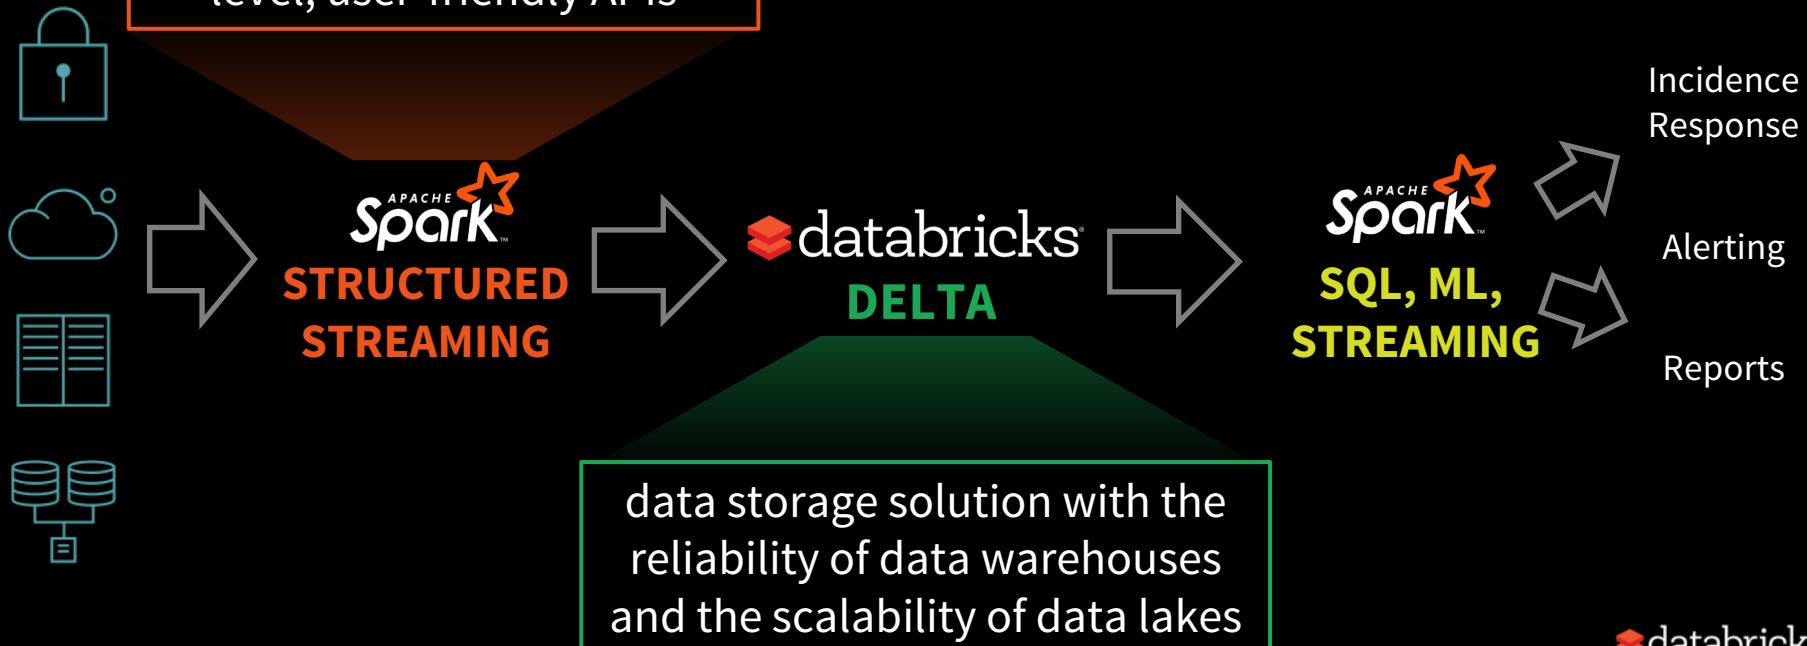
DATA



HADOOP STORAGE

# Data Pipelines with and DELTA

fast, scalable, fault-tolerant  
stream processing with high-  
level, user-friendly APIs



# More Info

## Structured Streaming Programming Guide

<http://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>

## Databricks blog posts for more focused discussions on streaming

<https://databricks.com/blog/category/engineering/streaming>

## Databricks Delta

<https://databricks.com/product/databricks-delta>



VIDEO ARCHIVE

SAIS 2018

SCHEDULE

TRAINING

VENUE + HOTEL

SPONSORS

ORGANIZED BY  
 databricks

REGISTER TODAY >



# SPARK+AI SUMMIT 2018

JUNE 4 - 6 | SAN FRANCISCO

ORGANIZED BY databricks

COUNTDOWN TO SUMMIT

47  
DAYS

07  
HOURS

26  
MINUTES

REGISTER TODAY >





# Thank you!

 @tathadas